

COSC 3360

Spring 2020

First Assignment

Jehan-François Pâris
jfparis@uh.edu



A very simple case


- **NCORES 1** // number of cores
- START 120** // new process at T=1200
- PID 23** // process ID
- CORE 100** // request CORE for 100 ms
- TTY 5000** // 5000 ms user interaction
- CORE 80** // request CORE for 80 ms
- SSD 1** // request SSD for 1 ms
- CORE 30** // request CORE for 30 ms
- SSD 1** // request SSD for 1 ms
- CORE 20** // request CORE for 20 ms



The model

We have


- One single-core CPU
 - NCORES = 1
- One SSD
- Many user windows
- Two CPU queues
 - Interactive
 - Non-interactive



Process states

A process can be

- **Running**
 - It occupies a core
- **Ready**
 - It waits for a core
- **Blocked**
 - It wait for an I/O completion





The solution (I)

- **NCORES 1**
 - CPU has one core
- **START 120**
 - Set clock at $T=120$ ms
- **PID 23**
 - Record arrival of process 23 at $T=120$ ms



The solution (II)

■ CORE 100

- $T=120\text{ms}$
- Allocate core to process 23 for 100ms
- Move 23 to RUNNING state
- Core completion at $T=120+100=220\text{ms}$

■ TTY 5000

- $T=220\text{ms}$
- Release core
- Move 23 to BLOCKED state
- TTY completion at $T=220+5,000=5,220\text{ms}$



The solution (III)

■ CORE 80

- $T=5,220\text{ms}$
- Allocate core to process 23 for 80ms
- Move 23 to RUNNING state
- Core completion at $T=5,220+80=5,300\text{ms}$

■ SSD 1

- $T=5,300\text{ms}$
- Release core
- Move 23 to BLOCKED state
- SSD completion at $T=5,300+1=5,301\text{ms}$



The solution (IV)

■ CORE 30

- $T = 5,301\text{ms}$
- Allocate core to process 23 for 30ms
- Move 23 to RUNNING state
- Core completion at $T=5,301+30=5,331\text{ms}$

■ SSD 1

- $T=5,331\text{ms}$
- Release core
- Move 23 to BLOCKED state
- SSD completion at $T=5,331+1=5,332\text{ms}$



The solution (V)

- **CORE 20**

- ☐ $T=5,332\text{ms}$
- ☐ Allocate core to process 23 for 20ms
- ☐ Move 23 to RUNNING state
- ☐ Core completion at $T=5,332+20=5,352\text{ms}$

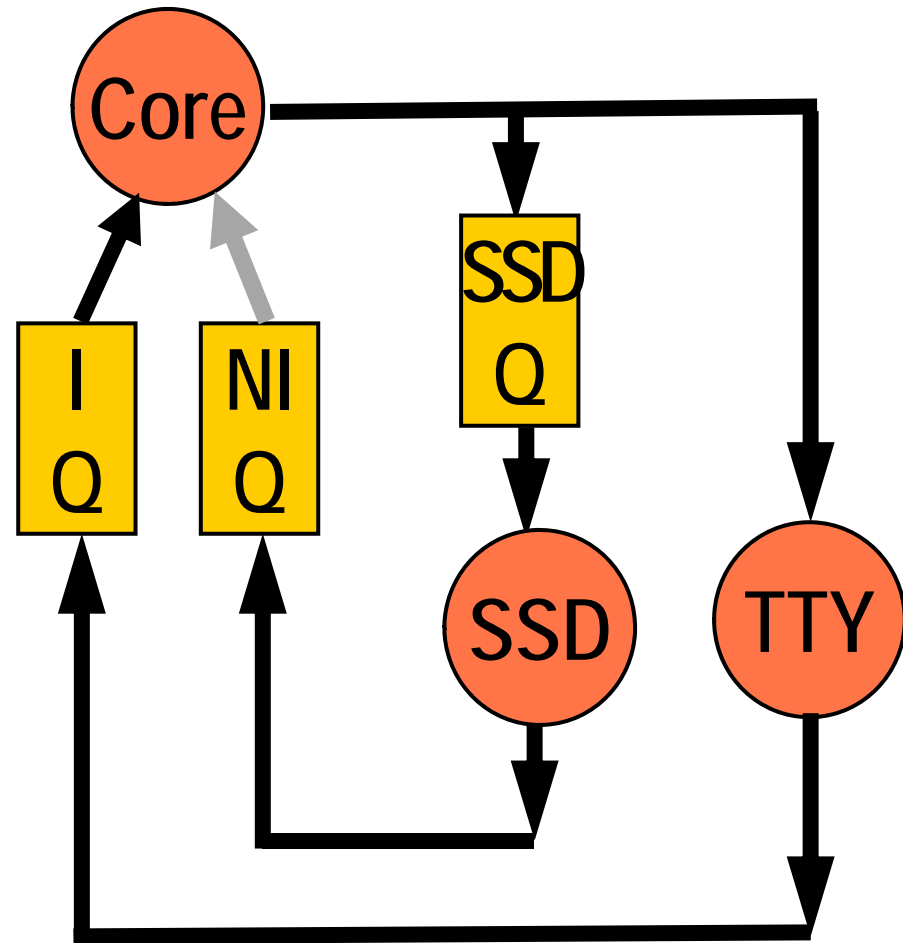
- ***Process terminates***

- ☐ Move 23 to TERMINATED state



Another way to look at it


NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20





T=120ms Process 23 arrives
Gets core until T = 220ms

NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20




[illegible]



T=5,220ms Process 23 gets core until = 5,300ms


NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20





T=5,300ms Process 23 releases core
Gets SSD until = 5,301ms


NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20





T=5,301ms Process releases SSD
Gets core until = 5,331ms


NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20





T=5,331ms Process releases core
Gets SSD until = 5,332ms


NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20





T=5,332ms Process releases SSD
Gets core until = 5,352ms


NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20





T=5,352ms Process terminates

NCORES 1
START 120
PID 23
CORE 100
TTY 5000
CORE 80
SSD 1
CORE 30
SSD 1
CORE 20





The completion events

- Completions are events
 - Mark end of a step
 - Notify it is time to move to next step
 - Next allocation step
- Scheduling is trivial when there is only one process
 - Not true otherwise



Handling two processes

NCORES 1

START 0

PID 0

CORE 10

Process 0

SSD 1

CORE 30

START 5

PID 1

CORE 20

Process 1

SSD 0

CPU 40



The solution (I)

Time	Command	Action(s)
0	NEW 0	Process 0 starts
0	CORE 10	P0 gets core until T=10ms
5	NEW 5	Process 1 starts
5	CORE 20	Core is busy: P1 enters NI queue P1 is in READY state
10	SSD 1	P0 releases the CPU Gets SSD until T = 11 ms P1 gets CPU until t = 30ms P1 is in RUNNING state



The solution (II)

Time	Command	Action(s)
11	CORE 30	<i>CPU is busy:</i> P0 enters NI queue P0 is in READY state
30	SSD 0	P1 releases the CPU Gets SSD until T=30 ms P0 gets CPU until T=60ms
30	CORE 40	<i>CPU is busy:</i> P1 enters NI queue P1 is in READY state



The solution (III)

Time	Command	Action(s)
60		P0 releases CPU and terminates P1 gets CPU until $T=100\text{ms}$
100		P1 releases CPU and terminates



Another way to look at it

NCORES 1

START 0

PID 0

CORE 10

SSD 1

CORE 30


START 5

PID 1

CORE 20

SSD 0


CPU 40





T = 0ms P0 gets core until T = 10ms

NCORES 1
START 0
PID 0
CORE 10
SSD 1
CORE 30
START 5
PID 1
CORE 20
SSD 0
CPU 40





T=5ms P1 waits for P0 to release core at T=30ms

NCORES 1

START 0

PID 0

CORE 10

SSD 1

CORE 30


START 5

PID 1

CORE 20

SSD 0

CPU 40





T=10ms

P0 gets SSD until T=11ms

P1 gets core until T=30ms

NCORES 1

START 0

PID 0

CORE 10

SSD 1

CORE 30


START 5

PID 1

CORE 20

SSD 0

CPU 40





T=11ms P0 waits for P1 to release core at T=30ms

NCORES 1

START 0

PID 0

CORE 10

SSD 1

CORE 30


START 5

PID 1

CORE 20

SSD 0

CPU 40





T=30ms

P1 gets SSD until T=30ms

P0 gets core until T=60ms

NCORES 1

START 0

PID 0

CORE 10

SSD 1

CORE 30


START 5

PID 1

CORE 20

SSD 0


CPU 40





T=30ms P1 waits for P0 to release core at T=60ms

NCORES 1
START 0
PID 0
CORE 10
SSD 1
CORE 30
START 5
PID 1
CORE 20
SSD 0
CPU 40





T=60ms

P0 to release core and terminates
P1 gets core until T=100ms

NCORES 1

START 0

PID 0

CORE 10

SSD 1

CORE 30

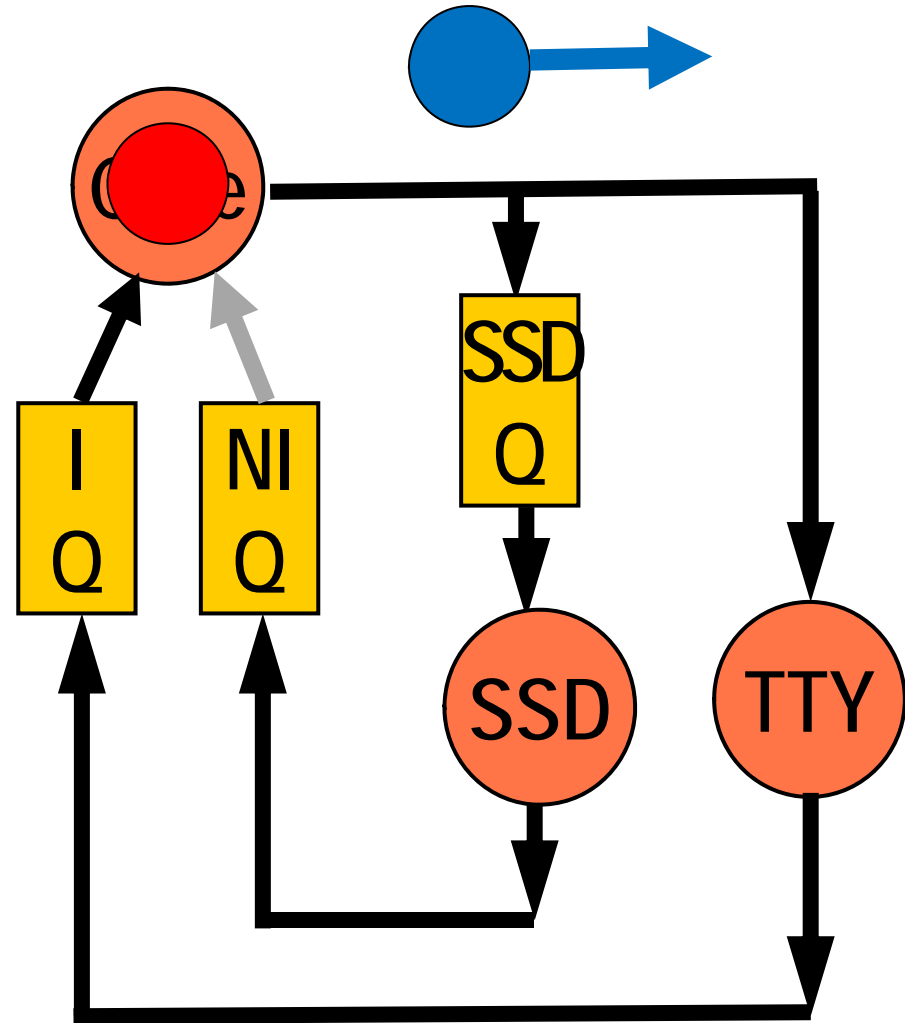
START 5

PID 1

CORE 20

SSD 0

CPU 40





START 0

CORE 10

SSD 1

CORE 30

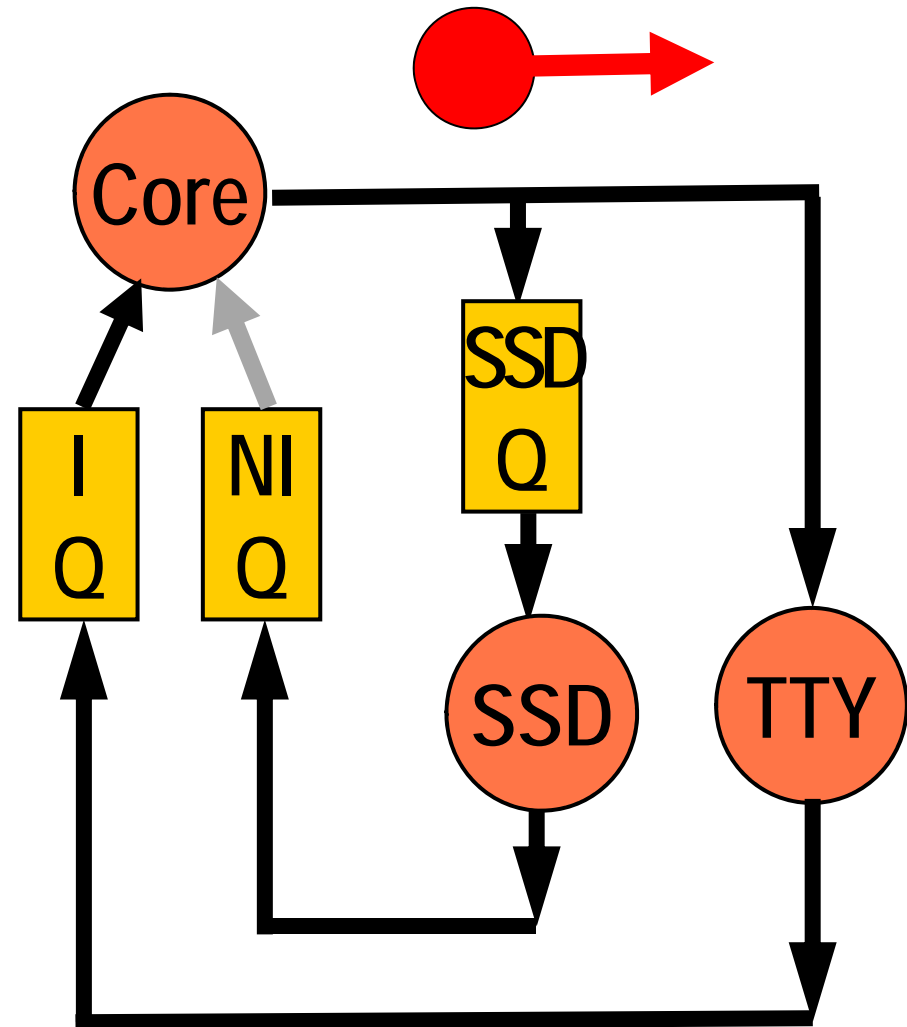
START 5

PID 1

CORE 20

SSD 0

CPU 40






Scheduling the CPU

- Have two queues
 - Interactive queue
 - Contains processes that have just completed a user interaction
 - Higher priority
 - Interactive queue
 - Contains all other processes
 - Lower priority


Example (I)

- ❑ When process 1 releases a core, process 4 will get it ahead of 2 and 3.



Example (II)

- ❑ If process 1 returns to READY state before 4 releases its core, it will get this core ahead of 2 and 3
- ❑ Otherwise process 2 will get the core





Handling parallel activities

- We only need to consider start times and completion times of each computational step
- Completion times are the most important
 - Release a device
 - Initiate the next request
 - Can be immediately satisfied if requested device is free
 - May require process to wait for device
 - Ready queue or disk queue



ENGINEERING THE SIMULATION



Simulating time

- ***Absolutely nothing happens to our model between two successive "events"***
- ***Events are***
 - Arrival of a new process
 - Completion of a computing step
 - Completion of a SSD access
 - Completion of a user interaction
- We associate an ***event routine*** with each event



Organizing our program (I)

- Most steps of simulation involve scheduling future completion events
- Associate with each completion event an event notice
 - **Time of event**
 - **Device**
 - **Process ID**
 - **Interactive/Non-interactive bit**



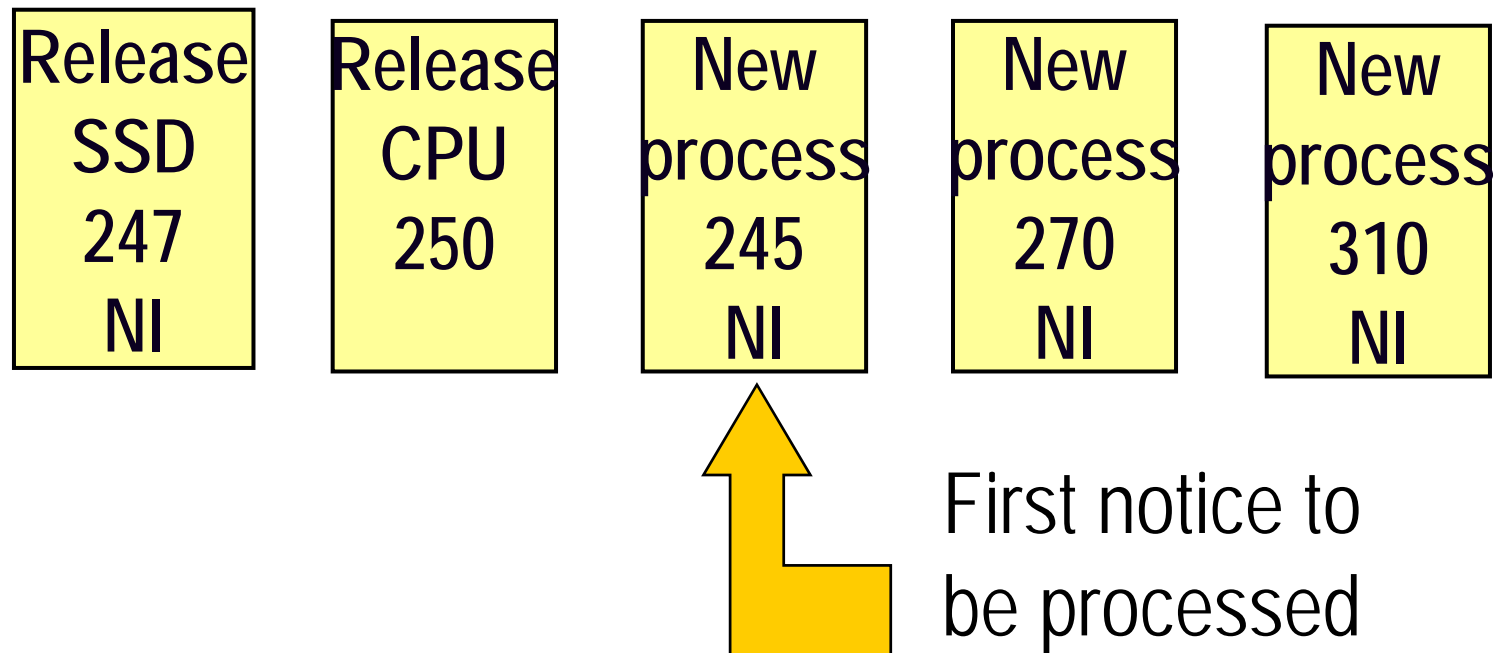
Organizing our program (II)

- Do the same with process starts
 - **Time of event**
 - **Process start**
 - **Process ID**
 - **Interactive/non-interactive bit**
 - **Set to non-interactive**



Organizing our program (III)

- Process all event notices in chronological order





Organizing our program (IV)

- Overall organization of main program

```
read in input file
schedule all process starts
while (event list is not empty) {
    process next event in list
} // while
print simulation results
```





Organizing our event list (I)

- As a priority queue
- Associating a completion time
 - With each core request
 - With each SSD request
 - With each user interaction
 - With the each new process arrival



Organizing our event list

- Process all event notices in chronological order



First notice to be processed
is at the head of the list



Arrival event routine

```
arrival(time, proc_id) {  
    process first request of new process  
} // arrival
```



Core request routine

```
core_request(how_long, proc_id, isinteract){
    if (nfreecores > 0) {
        nfreecores--;
        schedule completion at time
        current_time + how_long
        for process proc_id;
    } else {
        if (isinteract == interactive) {
            queue proc_id in i_queue
        } else {
            queue proc_id in ni_queue
        } // inner if
    } // outer if
} // core_request
```



Core completion routine

```
core_release (proc_id){
    if (i_queue is not empty) {
        pop first request in i_queue
        schedule its completion at
        current_time + how_long;
    } else if (ni_queue is not empty) {
        pop first request in ni_queue
        schedule its completion at
        current_time + how_long;
    } else {
        nfreecores++;
    } //if
    process next process request;
} // core_release
```



SSD request routine

```
ssd_request(how_long, proc_id){  
    if (ssd == FREE) {  
        ssd = BUSY;  
        schedule completion at time  
        current_time + how_long  
        for process proc_id;  
    } else {  
        queue process proc_id in  
        ssd queue;  
    } // if  
} // ssd_request
```




SSD completion routine

```
ssd_release (proc_id, &isinteract){  
    isinteract = non_interactive;  
    if (ssd queue is not empty) {  
        pop first request in ssd queue  
        schedule its completion at  
        current_time + how_long;  
    } else {  
        ssd = FREE;  
    } // if  
    process next process request;  
} // ssd_release
```



User request routine

```
user_request (how_long, proc_id){  
    schedule completion at time  
    current_time + how_long  
    for process process_id;  
} // user_request
```



User completion routine

```
user_release (proc_id, &isinteract){  
    isinteract = interactive;  
    process next process request;  
} // user_release
```



Overview (I)

- ***Input module***

- Schedules all ARRIVAL events

- ***Main loop***

- Pops next event from event list
 - ARRIVAL
 - CORE completion
 - SSD completion
 - TTY completion



Overview (II)

- ***ARRIVAL event***
 - Starts a core request



Overview (III)

- ***Core request***

- If a core is free

- Schedules a CORE completion event

- ***CORE completion event***

- May schedule a CORE completion event

- Starts next request

- SSD or TTY



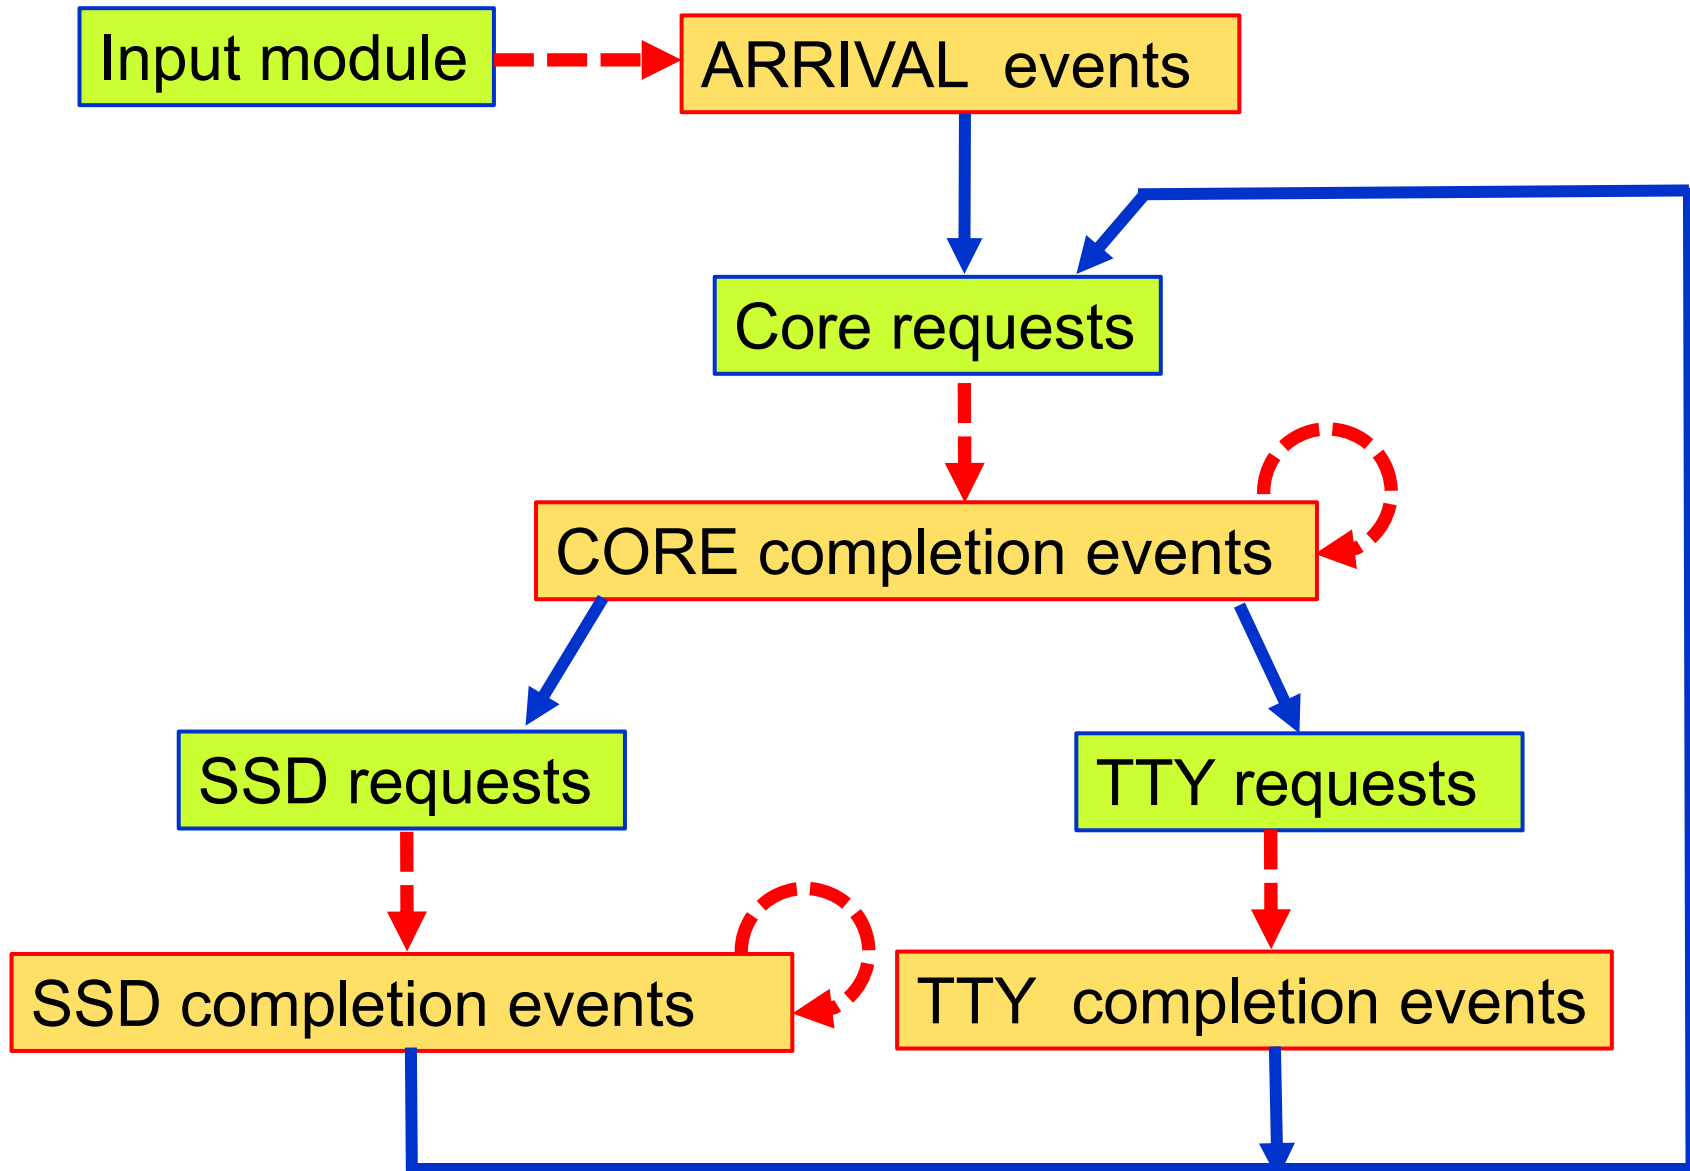
Overview (IV)

- ***SSD request***
 - If SSD is free
 - Schedules an SSD completion event
- ***SSD completion event***
 - May schedule a SSD completion event
 - Starts next request
 - CORE



Overview (V)

- ***TTY request***
 - Schedules an TTY completion event
- ***TTY completion event***
 - Starts next request
 - CORE





Explanations

- Green boxes represent conventional functions
- Amber boxes represent events and their associated functions
- Continuous blue arrows represent regular function calls
- Red dashed lines represent the scheduling of specific events



Finding the next event

- If you do not use a priority list for your events, you can find the next event to process by searching the **lowest value** in
 - The process start times in the process table
 - The completion times in the device table
 - The display completion timed in the process table



AN IMPLEMENTATION

- My main data structures would include:
 - An input table
 - A process table
 - A device table

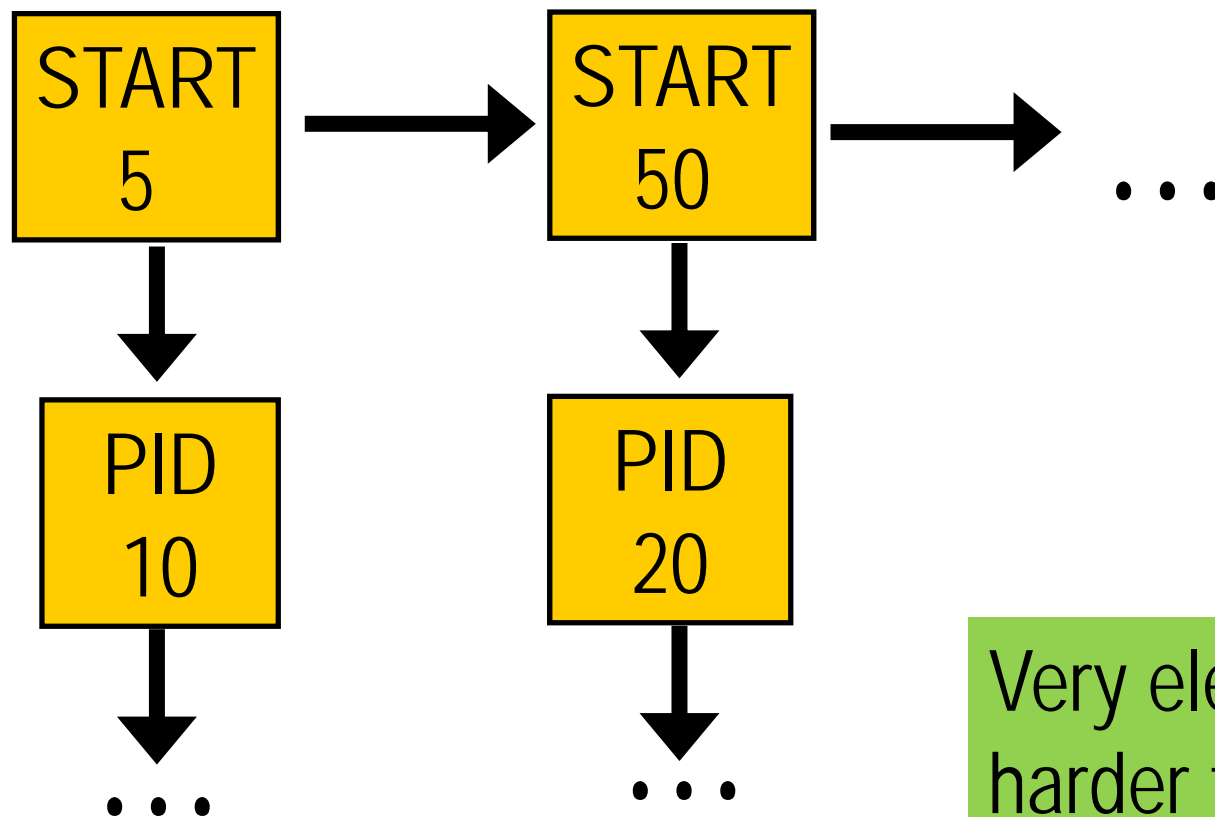


The input table

- Stores the input data
- Line indices are used in process table

Operation	Parameter
START	5
PID	10
CORE	20
SSD	0
CORE	20
START	50
PID	20
...	...

A full list implementation of the input table



Very elegant but
harder to debug



The process table (I)

PID	Start Time	First Line	Last Line	Current Line	State
10	5	0	4	varies	varies
20	50	5	
	



The process table (II)

- One line per process
 - ***Line index is process sequence number!***
- First column has start time of process
- First line, last line and current line respectively identify first line, last line and current line of the process in the input table
- Last column is for the current state of the process (READY, RUNNING or BLOCKED)



The device table (I)

Device	Status	Busy times
CPU	P0	15
SSD	-	-



The device table (II)

- One line per device
 - ***Line index identifies the device***
- First column has status of device
 - Number of free cores for CPU
 - Free/busy for SSD
- Last column is for the total of all busy times



READING YOUR INPUT

- You must use I/O redirection
 - `assign1 < input_file`
- ***Advantages***
 - Very flexible
 - Programmers write their code as if it was read from standard input
 - No need to mess with `fopen()`, `argc` and `argv`



Detecting the end of data

- The easiest ways to do it
- If you use `scanf()`
 - `scanf(...)` returns `0` once it reaches the end of data
 - `if(scanf(...))`
- If you use `cin`
 - `cin` returns `0` once it reaches the end of data
 - `while (cin >> keyword >> time)`